ಆದೇಶ AI

Aadesh AI

Visual Handbook · v3.3

AI drafts. Human verifies. Officer remains responsible.

A personal AI drafting assistant

that learns your style and writes formal orders in minutes.

By Srinivasa T, Founder

aadesh-ai.in

April 26, 2026

CONTENTS

Table of contents

AT A GLANCE

Aadesh AI in one page

Eight numbers that summarise the product as of April 26, 2026.

aadesh-ai.in live website	**1** pilot user (Banu)	**Pilot** pricing in validation	**~95%** time saved per order
89% / 98% live / lab quality	**6** compliance layers	**May 12** RAI deadline (target)	**Pilot** break-even pending data

AI drafts. Human verifies. Officer remains responsible.

Status note • Compliance layers 1–5 are code-implemented but pending production migration and smoke test as of April 26, 2026. See Section 10A for honest per-layer status.

Read the next 12 sections to see how each of these numbers is achieved, what is built, what is in progress, and what could go wrong.

SECTION 1

The basics

> *AI drafts. Human verifies. Officer remains responsible.*

Q1. What is Aadesh AI?

Aadesh AI is a self-service web product that helps officers and professionals draft formal documents using Artificial Intelligence. The user uploads twenty of their best past documents — the AI learns their personal drafting style — and from then on, every new case file becomes a complete formal draft in two minutes, written in their own style.

Aadesh does not train a shared AI model on your documents. Your reference documents are used only as private context for your own drafts.

Q2. What does the name "ಆದೇಶ" mean?

ಆದೇಶ (Aadesh) is the Kannada word for "Order." It is the same word used in formal Karnataka government documents — for example, "ಆದೇಶ ಸಂಖ್ಯೆ" (Order Number). The product is named after what it helps people write.

Q3. Who is the primary user, and who is the responsible authority?

These are two different people, and the distinction matters legally and commercially.

Primary user — caseworker / drafting assistant. The person who uploads case files, runs the AI, edits the draft. Our pilot user Banu is a DDLR caseworker.

Responsible authority — officer / signatory. The person who reads the final draft, applies independent reasoning, signs the order, and bears legal responsibility for it.

Aadesh AI assists the caseworker. The officer signs and remains accountable.

Q4. Who can use it?

Two broad groups: government officers (DDLR, Tahsildars, AC, DC, BBMP staff) and private professionals (advocates, chartered accountants, notaries, consultants). Anyone who writes formal documents repeatedly. The next page shows the full audience tree.

Q5. Why does this product exist?

Two reasons. First — money. Caseworkers and officers pay private human drafters between ₹1,000 and ₹2,000 per order, out of their own pocket. For someone writing thirty orders a month, that is roughly ₹45,000 per month coming out of their salary.

Second — time. Even when an officer drafts an order without a private drafter, it takes one to three hours per document. Aadesh AI reduces the cost dramatically and gives back most of the time.

Q6. Why now?

Three reasons that all point to this exact window.

AI has only recently become good enough at Sarakari Kannada to do this work properly. With Claude Sonnet 4.6 and prompt caching, output quality crosses ninety percent.

Karnataka has set up a Responsible AI Committee with an interim report expected within sixty days of formation (March 2026). Our internal target is to submit our representation before May 12, 2026, aligning with the committee's early consultation window.

Sarvam AI received ₹247 crore in government funding. The first-mover window is six to eighteen months.

Q7. How should I think about this product?

Two analogies that work together.

The junior clerk who never forgets. Imagine a new junior clerk joins your office. You hand them your best twenty past orders and say: "Read these. This is how we write here." They read every word. They never forget. From their first day, they draft new orders in your style. They work at any hour, never get tired, never ask for a raise. That is Aadesh AI.

The personal whiteboard. The AI is a blank whiteboard. You write on it by uploading your best documents. Whatever you write, the AI mimics. If you transfer to a new office, you wipe the whiteboard clean and start fresh.

AUDIENCE

Who Aadesh AI serves

Two broad audiences for the same product: government officers who write formal land-revenue or municipal orders, and private professionals who write formal documents repeatedly.

AADESH AI

↓ ↓

GOVERNMENT OFFICERS	PRIVATE PROFESSIONALS
• **DDLR caseworkers · pilot user: Banu**	• Advocates · legal notices
• Tahsildars	• Chartered Accountants · compliance letters
• Assistant Commissioners (AC)	• Notaries · affidavits
• Deputy Commissioners (DC)	• Consultants · formal reports
• BBMP staff	

Same product, different document types. Each user trains the AI on their own past documents — so each user gets their own private version of Aadesh tuned to their writing style.

SECTION 2

The problem we solve

Q1. What does an officer do today to write one order?

Seven steps. (1) Receives a forty-to-sixty page scanned case file. (2) Reads everything and picks out the facts that matter — one to two hours. (3) Cross-checks survey numbers and party names across pages, since they often disagree — thirty to sixty minutes. (4) Asks a senior officer about tricky points. (5) Hands the file to a private drafter who writes the Kannada order — one to two days, ₹1,000 to ₹2,000 in cash. (6) Sends the draft to a typist. (7) Reviews, signs, issues. Total: three to six hours of attention spread over two to three days, plus ₹1,000 to ₹2,000 in cash.

Q2. How big is the problem in numbers?

A typical caseworker writes between twenty and sixty orders a month. At an average of ₹1,500 per order × thirty orders = ₹45,000 per month, every month, from the caseworker's salary.

Cost comparison — drafter vs Aadesh AI

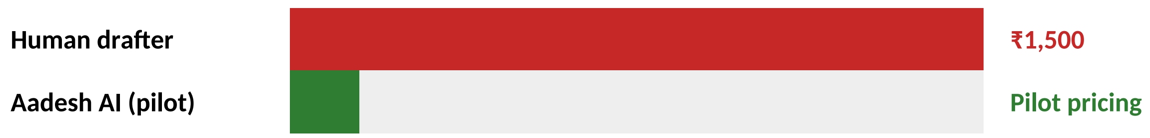

Aadesh AI is targeting a price point that delivers significant savings vs the human drafter while remaining sustainable for the platform. Final pricing is in pilot validation.

Comparison: today versus with Aadesh AI

Step in the workflow	Today (manual)	With Aadesh AI
Read the case file	20 minutes	0 minutes (AI reads it)
Extract key facts	15 minutes	0 minutes (AI extracts)
Draft the order	60 to 120 minutes	0 minutes (AI drafts)
Review and correct	30 minutes	5 to 10 minutes
Total caseworker / officer time	2 to 3 hours	5 to 10 minutes
Total cash cost	₹1,000 to ₹2,000	Pilot pricing in validation

Q3. Why don't officers just write everything themselves?

Many do. But it takes one to three hours per order. With twenty pending cases, that becomes the entire workday — and the queue keeps growing. The drafter exists because time is the bigger cost than money.

Q4. Why doesn't a generic ChatGPT solve this?

Three reasons. (1) Sarakari Kannada is a specialised formal dialect that generic AI does not handle reliably. (2) Every district has its own writing style — Bangalore Urban writes orders differently from Mysore, Gulbarga, Kolar. (3) Generic AI does not learn personal style — every prompt starts from zero.

Aadesh AI is built specifically for this combination of constraints.

SECTION 3

The user's journey

Q1. How does someone start using Aadesh AI?

One-time setup, ten minutes. (1) Sign up at aadesh-ai.in using Google login or email. (2) Choose office type and district. (3) Upload five to twenty of your own finalised, signed orders. (4) The AI reads them and learns your drafting style. (5) Done. Ready to use.

The setup-versus-daily-use flow

1. Sign up	→	2. Upload references	→	3. Daily use begins
2 minutes		one-time, 5–10 min		every order, 4 min

Q2. What does daily use look like?

Four minutes per order. (1) Click "New Order." (2) Upload the case PDF. (3) The AI reads the PDF and extracts the facts. (4) The Entity Lock screen appears: you verify the facts before anything is drafted. (5) The AI asks four to five clarifying questions. (6) The AI generates the full order in your style. (7) You review on screen and edit anything. (8) Download the .docx file. (9) Print, sign, submit.

Active caseworker time: about five minutes total.

Q3. What happens when I edit a draft?

Your edits are saved. Future orders learn from those edits — the AI adjusts its understanding of your style based on what you change. Quality improves with use.

Q4. What if I transfer to a new office?

Use Transfer Mode. Old data is archived. The AI clears. You upload your new office's reference orders and the AI re-learns. First two weeks free.

Q5. Where do my orders live? Can other people see them?

Your private space, completely isolated. No other officer — including officers in your own office — can see your orders. The database has row-level security: every read query is filtered by your user ID.

Users can request data export and deletion at any time. During the pilot, deletion requests are processed by the founder/admin within 48 hours and logged. Self-service deletion is planned before wider rollout.

SECTION 4

The PDF journey

Q1. What happens when I upload my case PDF?

The system processes scanned government PDFs including CamScanner files. Clear printed Kannada works best. Poor scans, rotated pages, handwriting, or damaged sections may require verification on the Entity Lock screen. From the readable parts, the AI extracts: party names, survey numbers, village/taluk/district, case number, dates, and what relief is being asked.

The four-step PDF flow

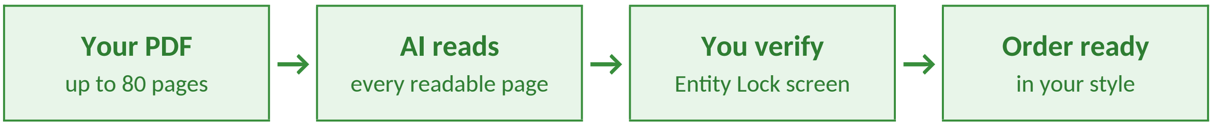

Q2. What if some details in the PDF disagree with each other?

This happens often. The survey number on page four may not match page twelve. The appellant's age in the appeal memo may differ from the Aadhaar card. The extent recorded in the tippani may differ from the registered deed.

The AI flags every disagreement on the Entity Lock screen. You decide which is correct before any draft is written. The AI follows a fixed precedence rule when sources disagree. You can override the AI's default choice.

Rule 18 conflict precedence pyramid

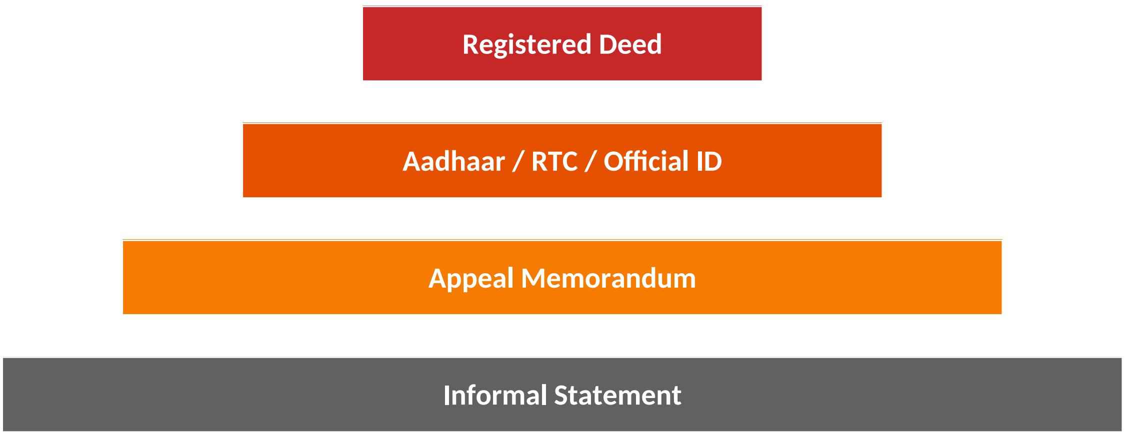

When two sources disagree, the higher tier wins. Officer can override on Entity Lock screen.

Q3. What if a fact is missing in the PDF?

The system is instructed not to fabricate. If a fact is missing or ambiguous, it leaves a placeholder like [___] in the draft and adds a note in the Review Note section. It will also ask you a clarifying question instead of guessing.

Q4. Are there limits on the PDF?

Up to eighty pages, up to thirty megabytes. PDF, DOCX, JPG, and PNG are all accepted.

SECTION 5

What makes Aadesh different

Q1. Can't a generic AI tool do this?

In our tests, generic tools were inconsistent for Sarakari Kannada legal drafting and did not preserve house-style conventions. Some tools showed weakness with Kannada inputs. None of them learn personal drafting style across sessions.

Generic AI tools do not provide Aadesh's combined workflow: per-user style learning, Entity Lock, officer reasoning, manifest hashing, Assistance Report, and Sarakari Kannada drafting in one product.

The per-district style problem

District	Style characteristic
Bangalore Urban	Formal urban phrasing, city-specific land terminology
Mysore	Different section structures, different fixed text blocks
Gulbarga	North Karnataka Kannada variations
Kolar	Distinct local terminology and format traditions
Shimoga	Unique Malnad-region administrative phrasing

No template works across all thirty-one Karnataka districts. Per-officer style learning is the moat: every officer trains the AI on their own files, so each user gets a different version tuned to their writing.

Q2. How does Aadesh actually learn each user's style?

Two methods working together.

Method 1 — Auto-prompt generator. When you upload twenty orders, the AI reads all of them and writes a custom set of instructions for itself, specific to your office. It detects your document structure, fixed text blocks, terminology, and decision patterns.

Method 2 — Reference library in working memory. Your twenty orders stay in the AI's context window during every generation. When drafting a new order, it has all twenty available to consult.

Q3. Will my style data ever be shared or used to train other people's AI?

Aadesh does not intentionally use user documents to train shared models, and processing is limited to generating the user's own draft. Other officers cannot see your data. If you delete your account, all your data is permanently removed.

SECTION 6

The technology behind it

Q1. Where does Aadesh AI live?

On a cloud server in Bangalore, India. The website is built with Next.js. The database is in Mumbai (Supabase, Indian region). The AI brain is Anthropic Claude Sonnet 4.6. Payments go through Razorpay. Domain is aadesh-ai.in.

The cloud architecture, in plain English

Officer browser aadesh-ai.in → **Bangalore VPS** Next.js app → **Mumbai database** Supabase → **AI brain** Claude Sonnet 4.6

Q2. What is the AI brain?

Claude Sonnet 4.6 from Anthropic, accessed via direct SDK with prompt caching. Best Kannada quality on the market today. Prompt caching lets us send a long, fixed system prompt once and reuse it cheaply for every order — bringing the marginal cost per order to roughly ₹20–30.

How Aadesh uses Claude's working memory

Three pieces of context go into the AI brain for every order. Two are cached for cost savings; one is unique to the case. Total ~78K tokens — under 8% of Claude's 1M-token context window.

Q3. How is data handled?

Where technically feasible, PII is masked before AI processing. Some document-reading steps may require secure AI processing of case material by the AI provider. Aadesh does not intentionally use user documents to train shared models, and processing is limited to generating the user's own draft.

Database in Mumbai. Server in Bangalore. Payments through Razorpay (India). Per-user isolation through database row-level security. Audit log of every access.

What happens inside the AI when you click Generate

Turn	What happens	Time	Cost
Turn 1 — Read PDF	AI processes each readable page and extracts identifiable facts	15–30 sec	~₹5
Turn 2 — Clarify	AI asks 4–5 specific questions; officer answers	60 sec	—
Turn 3 — Generate	AI writes full Kannada order in your style	30–60 sec	~₹14
Turn 4 — Self-audit	AI checks its own output; if a check fails, regenerates Turn 3 once with corrections	10–20 sec	~₹3

The 7 checks the AI runs on its own draft

CHECK 1

Sarakari Kannada tone

No transliteration, no English crutches, no AI-sounding sentences.

CHECK 2

Factual accuracy vs PDF

Every name, date, survey number, and extent matches the source PDF.

CHECK 3

Terminology table

Zero forbidden words. 39 forbidden / correct pairs are scanned.

CHECK 4

Fixed text blocks

Six boilerplate blocks (FT-1 to FT-6) match character-for-character.

CHECK 5

Structure

All 13 Appeal sections (or 8 Suo Motu sections) present in correct order.

CHECK 6

Allowed vs Dismissed logic

Dismissed orders correctly omit compliance directions.

CHECK 7

V3.2.2 rules

Rule 18 conflicts flagged. Rule 19 placeholders logged. Rule 20 advocate roll numbers stripped.

Q4. How does Aadesh check its own quality?

Self-validation. When you upload twenty orders, the system silently holds back three of them as test cases. It generates draft orders against those three and compares the AI output to your real signed orders. If the match score is below ninety percent, it asks you to upload more orders. Self-validation gives an automated quality score before any human sees a draft. Pilot feedback from real users (currently Banu) confirms whether that score correlates with real signing-ready quality.

Q5. Where is Aadesh built and processed?

The product, server (Bangalore VPS), database (Mumbai Supabase), and payments (Razorpay) are India-based. Some AI processing uses Anthropic under API and DPA-bound terms — this is the only step that touches a non-India server.

SECTION 7

The human stays in charge

> *Verification proves correctness. Reasoning proves authority.*

Q1. Who is legally responsible if a draft is wrong?

The officer who signs the order. Always. Aadesh AI is not an autonomous decision system. It is a human-in-the-loop drafting assistant for formal orders. The human verifies facts, applies reasoning, edits the draft, and signs the final order. The AI is a tool. The human is the author.

The six compliance layers (status: in progress, pending production smoke test)

Layer 6 · Digital Signature (Aadhaar eSign / DSC)
Planned — later phase

Layer 5 · Manifest Hash
In progress — code implemented, pending production verification

Layer 4 · Assistance Report PDF
In progress — code implemented, pending production migration / smoke test

Layer 3 · Audit Log (server-side)
In progress — implemented, pending production verification

Layer 2 · Officer Reasoning Block
In progress — code implemented, pending production smoke test

Layer 1 · Entity Lock (server-side enforced)
In progress — code implemented, pending production migration / smoke test

Server-side enforcement: The system is designed to block generation unless upload, Entity Lock, officer reasoning, and manifest seed exist in the server-side state record. Direct API access without these is intended to be unable to produce a draft.

Q2. What does the Officer Reasoning Block capture?

Four short items the officer fills in before finalizing.

Item	What the officer captures
1. Key issue	What the case turns on (e.g., paper-only phodi without notice)
2. Evidence relied upon	Which documents and findings drove the decision

Item	What the officer captures
3. Conflict resolution	Why the officer accepted one source over another (Rule 18)
4. Why decision is justified	The officer's independent reasoning for the outcome

Q3. What is in the Assistance Report PDF that comes with every order?

A receipt showing the full chain of custody. After production compliance deployment, every generated order will include this report; today, the implementation is in pilot. The Assistance Report is an audit artifact, not a government-issued certificate.

Field	What it contains
1. Input PDF SHA-256 hash	Fingerprint of the case file the officer uploaded
2. Generated order content hash	Fingerprint of what the AI produced
3. Officer attestation hash	Fingerprint of the officer's typed-name attestation
4. Combined manifest hash	Single fingerprint that ties all of the above together
5. Timestamp	When the order was generated
6. Source file details	PDF size, page count, type
7. Extracted fields	What the AI read out of the PDF
8. Confirmed fields	What the officer confirmed on Entity Lock
9. Conflicts resolved	Each conflict and the officer's chosen value
10. Model / prompt metadata	V3.2.6, Claude Sonnet 4.6, etc.
11. Officer responsibility statement + footer	"AI-assisted draft verified by officer"

The chain of hashes — how tamper-proofing works

Input hash PDF SHA-256 → **Output hash** order content → **Attestation hash** officer name → **Manifest hash** combines all 3 → **Optional eSign** Aadhaar / DSC

Q4. What is the BNS forgery rule?

Bharatiya Nyaya Sanhita is India's criminal code. If a wrong survey number ends up in a signed government document, that document becomes a "forged electronic record" under the BNS. The Entity Lock — the screen where the system is designed to require the officer to verify facts before drafting starts — creates evidence that the responsible human verified the facts before finalization.

This helps establish that Aadesh AI acted as a drafting assistant and that the officer/signatory retained responsibility for the final document.

Q5. What about the Karnataka RAI Committee?

Karnataka set up a Responsible AI Committee chaired by Kris Gopalakrishnan in March 2026. The committee is publishing risk classifications for AI used in government work. Our internal target is to submit our representation before May 12, 2026, aligning with the committee's early consultation window. We aim to be classified as medium-risk human-in-the-loop, not high-risk autonomous AI.

Q6. What about the Gujarat High Court ruling?

Recent Gujarat High Court AI policy shows that courts are highly sensitive to AI being used for legal reasoning, adjudication, or final order drafting. Aadesh therefore positions AI only as a drafting assistant and requires human verification, reasoning, editing, and signing. We treat this as a warning signal and design accordingly.

Q7. Is my citizen data treated well under DPDP?

Aadesh is designed to support DPDP-aligned controls: user consent, purpose-limited processing, access controls, audit logging, and deletion/export workflows. Some controls, including full retention/deletion policy and formal data-processing documentation, are being finalised before wider rollout. Formal retention period is being finalised before wider rollout.

SECTION 7A

What Aadesh does NOT do

Five explicit boundaries. These are not features that "could be added later" — they are intentional limits, required for the human-in-the-loop classification and for legal safety.

✕	**Does not auto-sign orders.**
✕	**Does not auto-submit orders.**
✕	**Does not replace statutory authority.**
✕	**Does not decide the case.**
✕	**Does not remove the officer's duty to verify facts, apply reasoning, edit, and sign.**

> *Aadesh AI is not an autonomous decision system. It is a human-in-the-loop drafting assistant for formal orders.*

SECTION 7B

Court / RAI evidence package

For every generated order, Aadesh can produce the following package on demand. This is the system's strongest defence in any legal or regulatory challenge. The Assistance Report is an audit artifact, not a government-issued certificate.

#	Evidence item	What it proves
1	Input PDF hash	The exact case file the officer uploaded was X
2	Entity Lock attestation	The officer typed their name and confirmed extracted facts
3	Officer Reasoning record	The officer applied independent reasoning, in their own words
4	Output hash	The exact draft the AI produced was Y
5	Manifest hash	Inputs, outputs, and attestation are tied together — no swap is possible
6	Assistance Report PDF	A human-readable receipt for any auditor
7	Prompt / model version	Which version of the system produced the draft (V3.2.6, Sonnet 4.6)
8	Timestamped audit log	Server-side proof of the sequence of events

> *Verification proves correctness. Reasoning proves authority.*

SECTION 8

The business

Q1. How much does it cost the user?

Aadesh AI uses prepaid credit packs. One completed order consumes one credit. Failed generations do not consume credits. Final pack pricing may change during pilot based on real API cost, usage volume, case complexity, and officer feedback.

Pricing approach during pilot

Aadesh AI is currently in pilot. Pricing is being validated against three real-world variables: (1) actual API cost per generation, which depends on case complexity and prompt-cache hit rate; (2) typical usage volume per officer per month; (3) feedback from pilot users on willingness to pay. Final published pricing will be confirmed after the first ten paying users.

For comparison, the human drafter that Aadesh AI replaces typically charges ₹1,000–₹2,000 per order. Aadesh AI is targeting a price point that delivers meaningful savings to the officer while remaining sustainable for the platform.

Q2. How does Aadesh make money?

Each completed order generates revenue net of the underlying AI processing cost, payment processing fee, and a share of fixed monthly infrastructure cost. The business is designed for sustainable per-order margins once usage patterns stabilise during pilot.

Q3. When will the business break even?

Break-even depends on final pricing, real per-order cost (measured during pilot), and user growth rate. A separate financial model maintained internally projects break-even at low double-digit active paying users, contingent on production cost staying within the budgeted range.

Year-1 trajectory (illustrative)

A typical SaaS trajectory in this category shows monthly losses for the first 3–4 months as user acquisition outpaces fixed-cost coverage, followed by break-even and modest profit in the second half of year one. Aadesh AI's actual trajectory will depend on pilot pricing decisions and user growth — both being validated now.

Q4. How will it scale beyond Karnataka land records?

Two-track go-to-market. Track 1 (now) is individual subscription. Track 2 (later, once 50+ users exist) is department procurement.

Aspect	Track 1 — Individual (now)	Track 2 — Department (later)
Who pays	Officer / caseworker from own pocket	Department from office budget
Approval	NONE needed	DC sufficient if under ₹5 lakh
Price	Pilot pricing in validation	TBD based on department size
Growth	Word of mouth between officers	Formal procurement process
Comparison	Like Netflix or ChatGPT	Like buying office software
Timeline	Live now (pilot)	Once 50+ users exist

Q5. What is the total Karnataka opportunity?

Roughly fifteen hundred potential users across land revenue, urban bodies, and other government departments. Total addressable market estimates are illustrative and contingent on final pricing decisions made during pilot.

SECTION 9

How it spreads

Q1. How does Aadesh grow without ad spend?

Three viral mechanisms built into the product itself.

The viral loop

Officer uses Aadesh finishes 5 orders fast → **DOCX has metadata** "Drafted with Aadesh AI" → **Colleague sees** on another officer's desk → **Colleague signs up** 3 free orders, hooked

Q2. What are the three mechanisms in detail?

Mechanism 1 — Drafts include an AI-assistance footer or metadata. Whether this footer remains on the final signed order is configurable by office policy.

Mechanism 2 — WhatsApp Referral. The officer shares a referral link in the district's officer WhatsApp group. The colleague gets three free orders. The officer earns five bonus credits.

Mechanism 3 — Peer Demo. Banu finishes five orders before morning tea. His colleague asks "How?" Banu shows his phone for sixty seconds. The colleague signs up at lunch.

Q3. Why don't paid ads work for this product?

Government officers do not search Google for "AI order drafting" — they do not know it exists. They do not trust ads for professional tools. They block cold calls. They trust other officers. Free word-of-mouth aligns with this preference.

SECTION 10

Where we are today

Q1. Is the product live?

Yes — aadesh-ai.in serves traffic. Health check passes. PM2 is stable. Supabase database is active. The end-to-end pipeline (sign-up, upload, generate, download) works for the pilot user.

Today's snapshot

Metric	Status
Live website	aadesh-ai.in — yes, live
Registered users	6
Test orders generated	29
Paying customers	0 (pilot phase)
Active pilot user	Banu, DDLR Bangalore South
System prompt version	V3.2.6 / V3.2.7 (~18,800 characters)
Quality benchmark — best lab test	98 / 100
Quality benchmark — live app	78 to 89 / 100
Compliance layers — production-ready	0 of 6 (5 in progress, 1 planned)
Days to internal RAI submission target	16

Quality benchmark journey

The journey is not monotonic. AnythingLLM-based approaches failed at 17–36%. Direct Claude SDK with V3.2.6 prompt is at 89% live. Best lab configuration crossed 98%.

Q2. What is the team?

Three actors. Srinivasa T — founder, non-technical, domain expert. Banu — first user, validator. Claude — CTO co-pilot, writes code, runs tests, reviews architecture. Lean by design. No hires planned until twenty active paying users.

Q3. What is the immediate priority?

Submit the Karnataka RAI Committee representation before our internal target of May 12. Apply production migrations and pass smoke tests on Layers 1–5. Get Banu signing his first real, billable order in the next thirty days.

SECTION 10A

Built vs Pending — honest status

Compliance and infrastructure status as of April 26, 2026. Items marked "in progress" have code implemented but pending production migration / smoke test. Do not treat any of these as "live in production" until the smoke test passes.

Item	Status
Entity Lock — server-side	In progress — code implemented, pending production migration / smoke test
Officer Reasoning Block	In progress — code implemented, pending production smoke test
Audit Log — server-side	In progress — implemented, pending production verification
Assistance Report PDF	In progress — code implemented, pending production migration / smoke test
Manifest Hash	In progress — code implemented, pending production verification
Digital Signature — Aadhaar eSign / DSC	Planned — later phase
One-credit-per-completed-order rule	In progress — code implemented, pending production billing smoke test
Self-service delete	Planned — pilot handled manually within 48 hours
PII masking before generation step	Partial — under review

After production smoke test passes, layers 1–5 and the one-credit rule will be re-classified as "Built — live in production." Until then, all claims about live behaviour reference the pilot environment only.

SECTION 11

The honest risks

Q1. What could go wrong on the technology side?

Quality of the Kannada draft. Hashes prove nothing if the draft itself is wrong. The biggest real test is whether the officer happily signs the AI draft without rewriting half of it.

Single AI vendor dependency. Aadesh depends on Anthropic. If Anthropic goes down, changes pricing, or restricts Indian traffic, the system is exposed. Mitigation: keep an OpenRouter and Sarvam fallback warm.

PDF processing memory. The current 1 GB virtual server can crash if multiple officers upload large PDFs simultaneously. Must upgrade to 4 GB by user number four.

Hosting empanelment gap. DigitalOcean is fine for individuals. Government procurement (Track 2) requires AWS Mumbai or Azure Pune migration.

Q2. What could go wrong on the business side?

Single-validator dependency. The pilot has one active validator. If they transfer offices, lose interest, or get too busy, validation stops. Mitigation: introduce one peer this month, identify a second user independently, record the demo and testimonial early.

Zero paying users today. Six testers and one active pilot is a strong signal — not proof. Three paying users in three months is the test of whether the willingness-to-pay hypothesis is right.

Sarvam AI catches up. They have ₹247 crore in government funding. The first-mover window is six to eighteen months — and counting.

Karnataka RAI mis-classifies the product. If classified as high-risk autonomous AI, the product may need to pause. Mitigation: a clear written submission emphasising human-in-the-loop architecture before May 12.

Q3. What is the worst case?

May 12 arrives with critical compliance gaps unfilled. The RAI committee classifies the product as high-risk autonomous AI. The product pauses in Karnataka. Aadhaar eSign and procurement-grade hosting must be added before re-submission. Track 2 delays by six to twelve months. Painful but survivable.

Q4. What is the best case?

The pilot officer signs the AI draft happily on the first real case. Word spreads in the district WhatsApp group. By month six, a viable cohort of paying users. By month twelve, the business is sustainable on Track 1 alone. Track 2 conversations with the Revenue Department begin in month nine. By year two, hundreds of users — final scale depends on pricing decisions and market validation.

SECTION 11A

Known limitations

Things the product currently cannot do, or does imperfectly. Knowing these matters more than not knowing them.

!	Poor scans may require manual correction on Entity Lock screen.
!	Handwriting recognition may be imperfect.
!	AI output must be reviewed by the officer before signing.
!	Legal citations and case-specific facts must be verified independently.
!	The product is currently in pilot — production smoke tests are pending.
!	Government procurement (Track 2) requires future hosting and compliance changes (MeitY-empanelled cloud).
!	Individual user adoption is being validated — willingness to pay is not yet proven beyond the pilot.
!	Self-service deletion UI is not yet built; deletion requests are manual during pilot.
!	Karnataka RAI classification is pending the committee's interim report.
!	Kannada legal responsibility wording is under final legal/Kannada review.

SECTION 12

Glossary

Every term in this document, explained in plain English.

Term	Plain meaning
ಆದೇಶ	Kannada for "Order." Used in formal Karnataka government documents.
Aadesh AI	The product. Personal AI drafting assistant for formal documents.
DDLR	District Deputy Land Records officer. The role our pilot user serves.
Tahsildar	Block-level revenue officer in Karnataka land administration.
AC / DC	Assistant Commissioner / Deputy Commissioner — district-level revenue officers.
BBMP	Bruhat Bengaluru Mahanagara Palike — Bangalore city corporation.
Sarakari Kannada	The formal Kannada legal/government dialect used in official orders.
Caseworker / drafter	Primary user of Aadesh. The person who drafts. (e.g., Banu)
Officer / signatory	Responsible authority. The person who signs and bears legal responsibility.
Entity Lock	Unskippable screen that forces the human to verify facts before AI drafts.
Officer Reasoning Block	A short capture of why the officer decided what they decided.
Manifest Hash	Cryptographic fingerprint that ties input, output, and attestation together.
Assistance Report PDF	A receipt attached to every order showing AI vs human contributions.
RAI Committee	Karnataka Responsible AI Committee — classifies AI products by risk level.
BNS	Bharatiya Nyaya Sanhita — successor to the Indian Penal Code.
DPDP Act	Digital Personal Data Protection Act 2023 — Indian privacy law.
Human-in-the-loop	AI system where a human reviews and signs the final output. Our target classification.
MeitY-empanelled	Government-of-India approved cloud hosting. Required for Track 2 procurement.
Prompt caching	Anthropic technology that lets a long fixed system prompt be reused cheaply.
Vision model	AI that "sees" PDF pages and extracts text from images, including scans.
PII	Personally Identifiable Information — names, addresses, survey numbers, etc.
Track 1 / Track 2	Individual subscription / Department procurement — our two-stage go-to-market.
Auto-prompt generator	AI that writes the AI's instructions, customised for each officer.

Term	Plain meaning
Self-validation	System holds back three of officer's files and grades AI output against them.
Sarvam AI	Indian-funded LLM company (₹247 crore from IndiaAI Mission). Reclassified competitor.
Razorpay	Indian payment gateway used for credit-pack purchases.
Supabase	Database and authentication platform we use, hosted in Mumbai region.
Anthropic	American AI company. Maker of Claude, our primary AI brain.
DPA	Data Processing Agreement — legal contract between data controller and processor.